

A Framework for Applying Neural Networks to Eddy Covariance Data

Dr. June Skeeter

june.skeeter@ubc.ca

University British Columbia

Eddy Covariance

Semi-continuous, ecosystem-scale measurements of energy, water, and trace gas fluxes.

- Noisy, voluminous data sets
 - Frequent gaps
 - Observational bias
- Well suited for machine learning!



**Burns Bog EC Station
Delta, BC**

Neural Networks

Universal approximators: can map any continuous function to an arbitrary degree accuracy.

- With enough hidden nodes, will fit **any** pattern in a dataset
 - Care must be taken to ensure the patterns are real
 - Early stopping allows
- Well suited for non-linear, multi-variate response functions
 - Capable **interpolation** and **extrapolation**

Commonly Cited Limitations

Issue	Solutions
Over-fitting	<ul style="list-style-type: none">- Model ensembles- 3-way cross validation- Pruning inputs
Black boxes models	<ul style="list-style-type: none">- Plot partial derivatives- Feature importance
Computationally expensive	<ul style="list-style-type: none">- Tensorflow- GPU processing

Objective

Provide a framework for applying NN models to EC data for descriptive analysis and inferential modelling.

- The [github repository](#) linked to this presentation has functional examples that can be used to apply NN models.
 - Runs in Python and Tensorflow
 - *GPU support not required*

Example Data

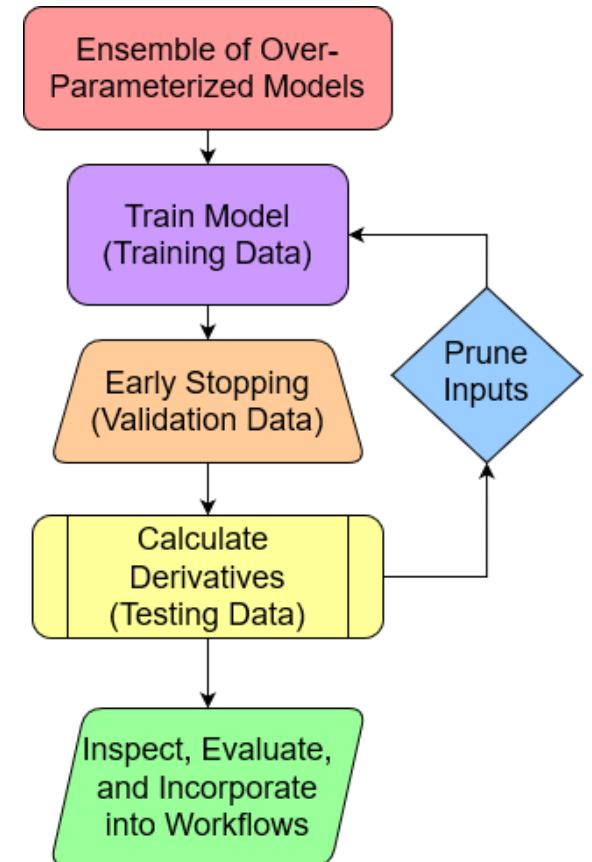
Burns Bog EC station

- Harvested peatland undergoing active restoration
- 8+ years of meteorological & flux (CO_2 and CH_4) data



Training Procedures

- Larger ensemble = more robust model
 - $N \leq 10$ for data exploration/pruning
- Three way cross-validation
 - Train/validate/test
- Early Stopping: after e epochs
 - $e = 2$ for pruning stage

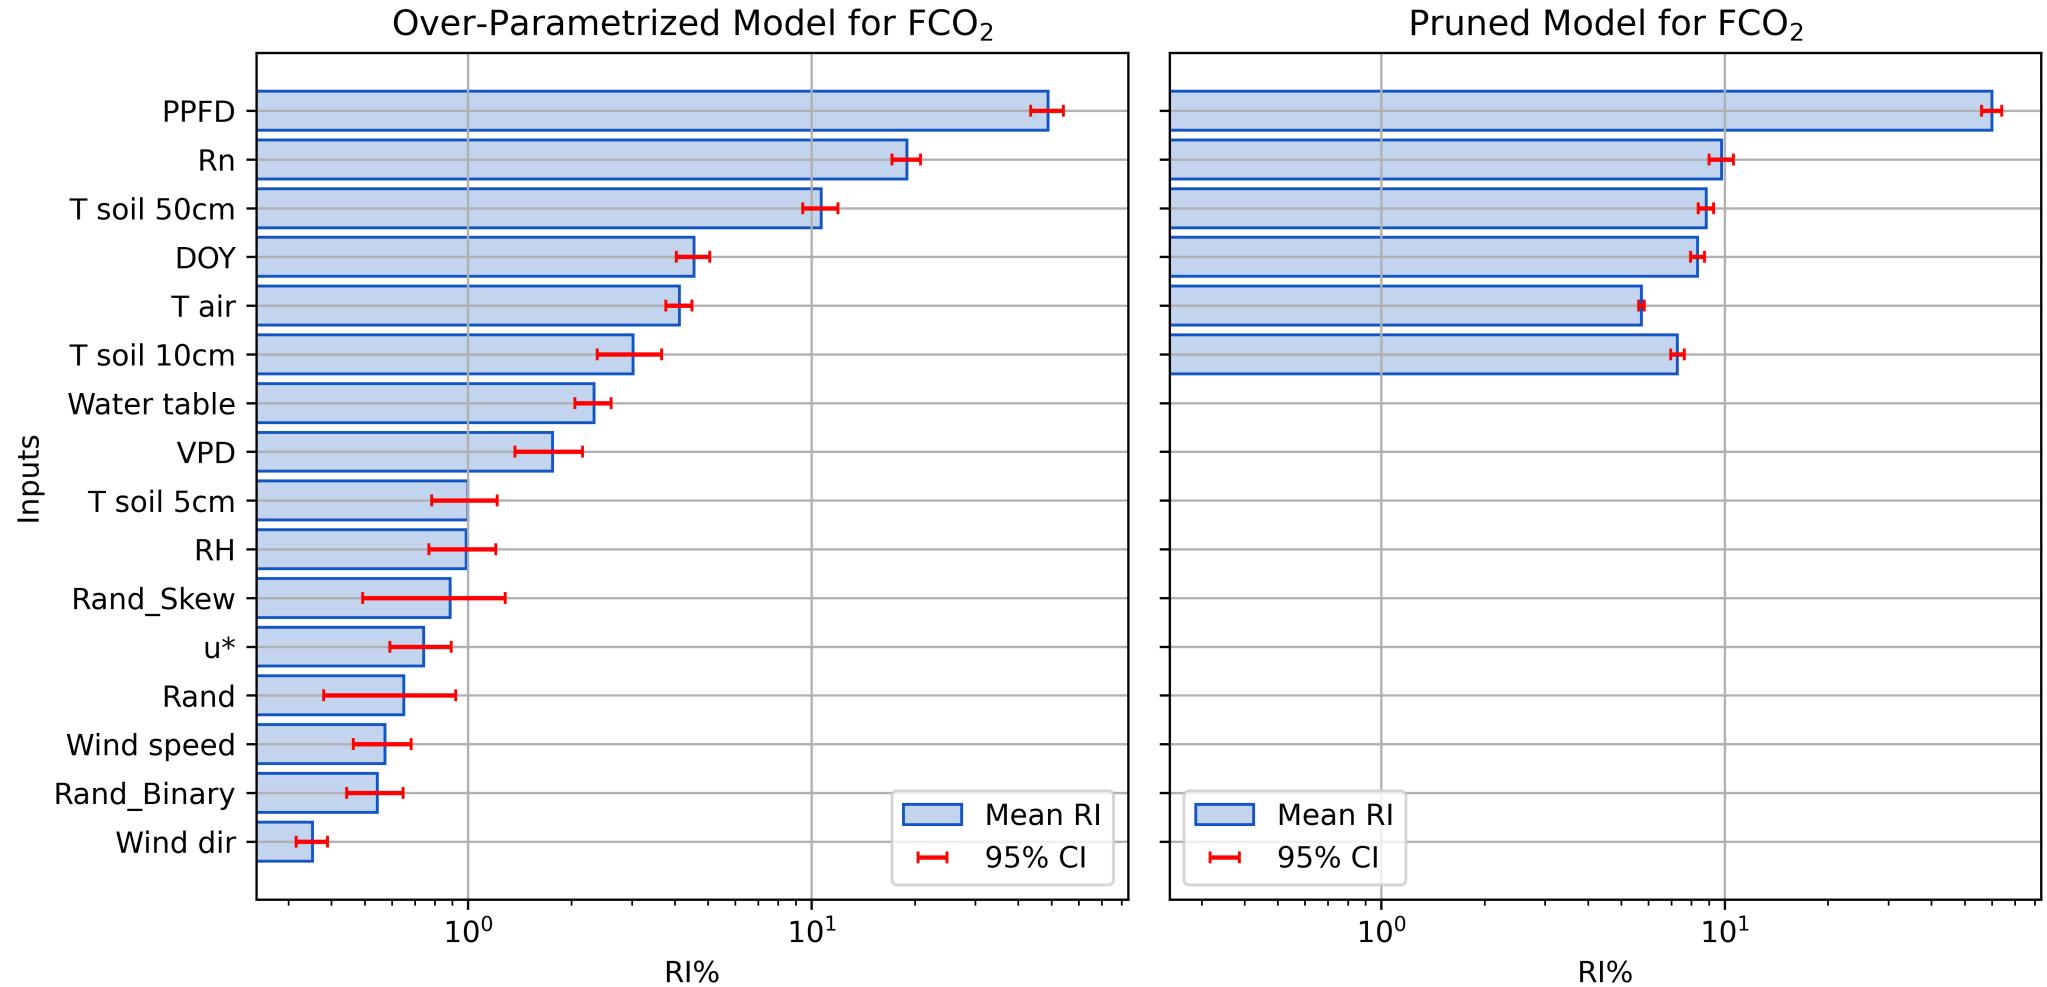


Pruning Inputs

Calculate partial first derivative of the output with respect to each input over test data domain.

- **Relative Influence (RI) of inputs**
 - Normalized sum of squared derivatives (SSD)
- Iteratively remove inputs with RI below a threshold
 - Use set of random scalars inputs to determine threshold
 - e.g., a float [0-1], a skewed float $[0-1]^{.25}$, and a binary integer (0/1)

Before and After Pruning FCO₂



The Final Model

Once pruning is complete, re-train the final production level model, excluding the random scalars

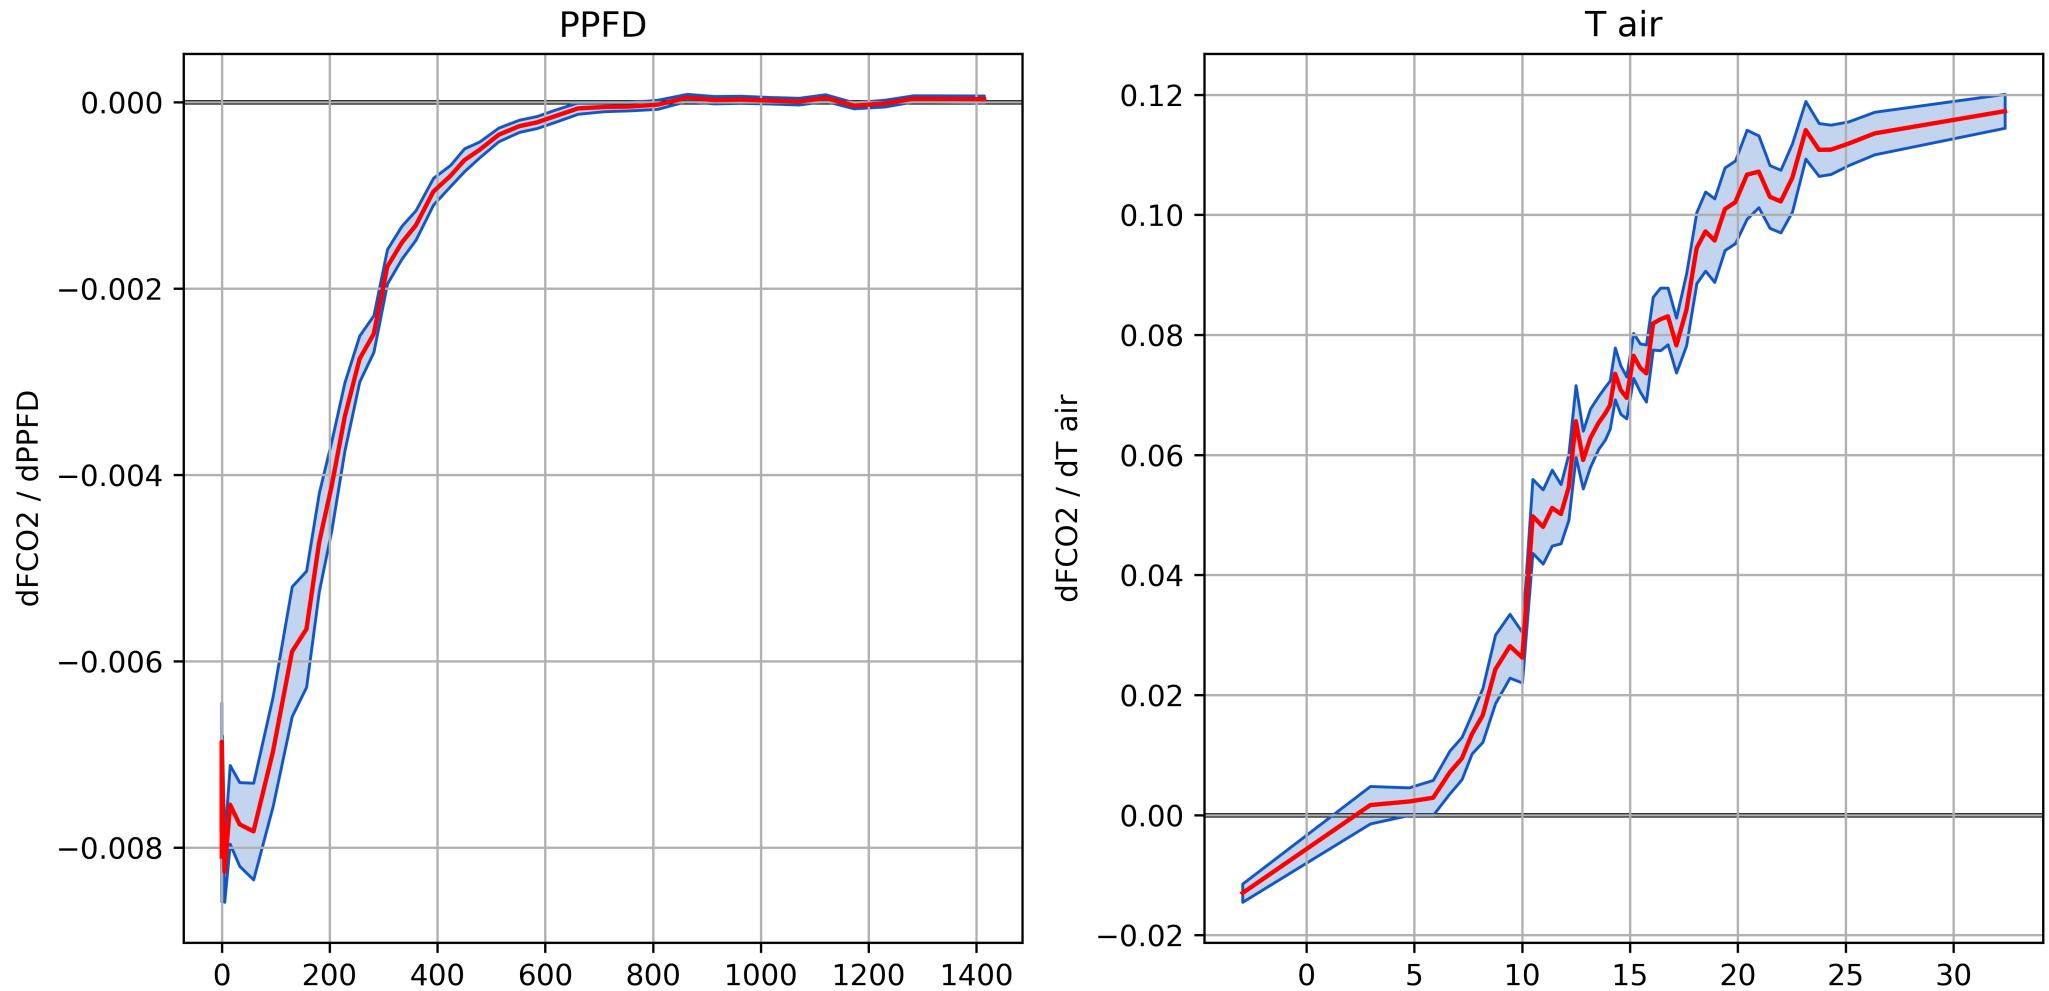
- Increase the ensemble size (e.g., $N = 30$)
 - Could increase early stopping criteria (e.g., $e = 10$)
 - Larger e drastically increases training
- Plot the model derivatives as a final check
 - If derivatives look implausible
 - Adjust inputs/parameters and try again

Plotting Derivatives

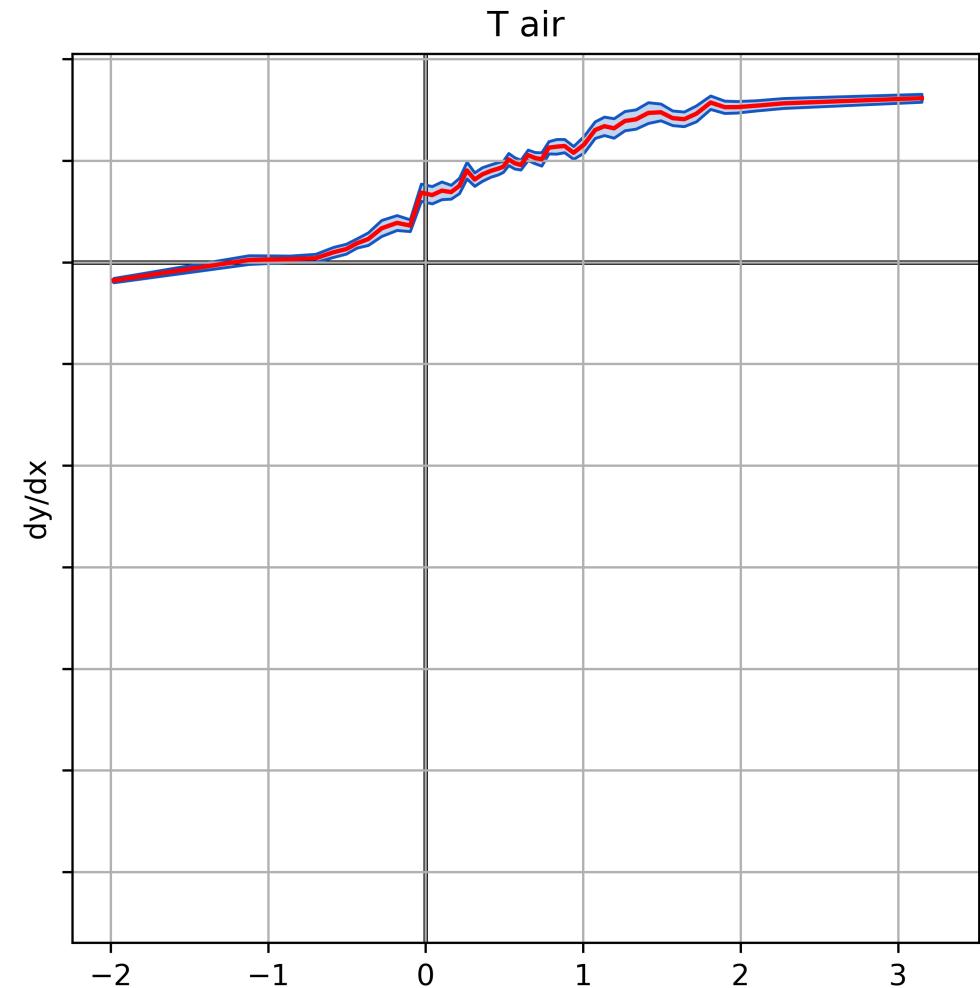
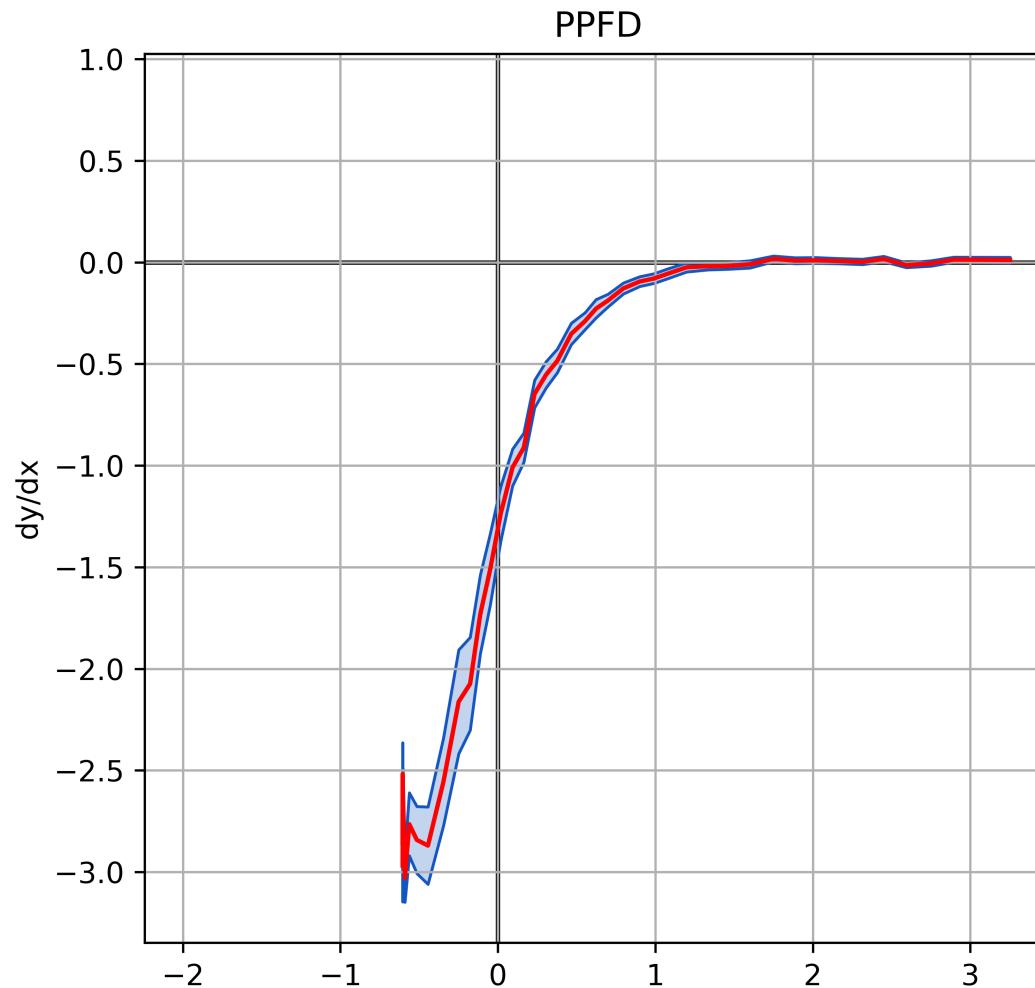
Helps ensure model responses are physically plausible

- An essential step and key advantage of NN models
- Raw derivatives show true feature responses
- Normalized derivatives scaled by input variance
 - Relative input effects on common scale
 - What the model “sees”
- 95% confidence intervals around derivatives indicate modeled confidence in relationships

Partial Derivatives of FCO₂



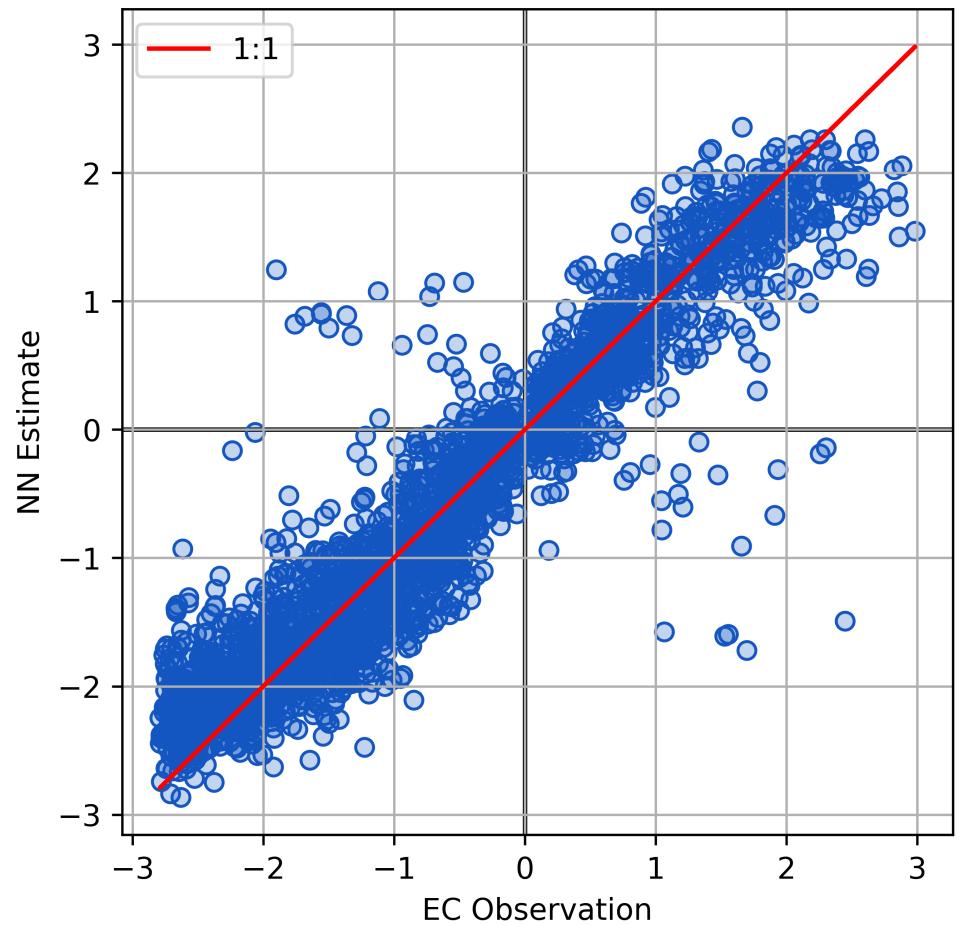
Normalized Derivatives of FCO₂



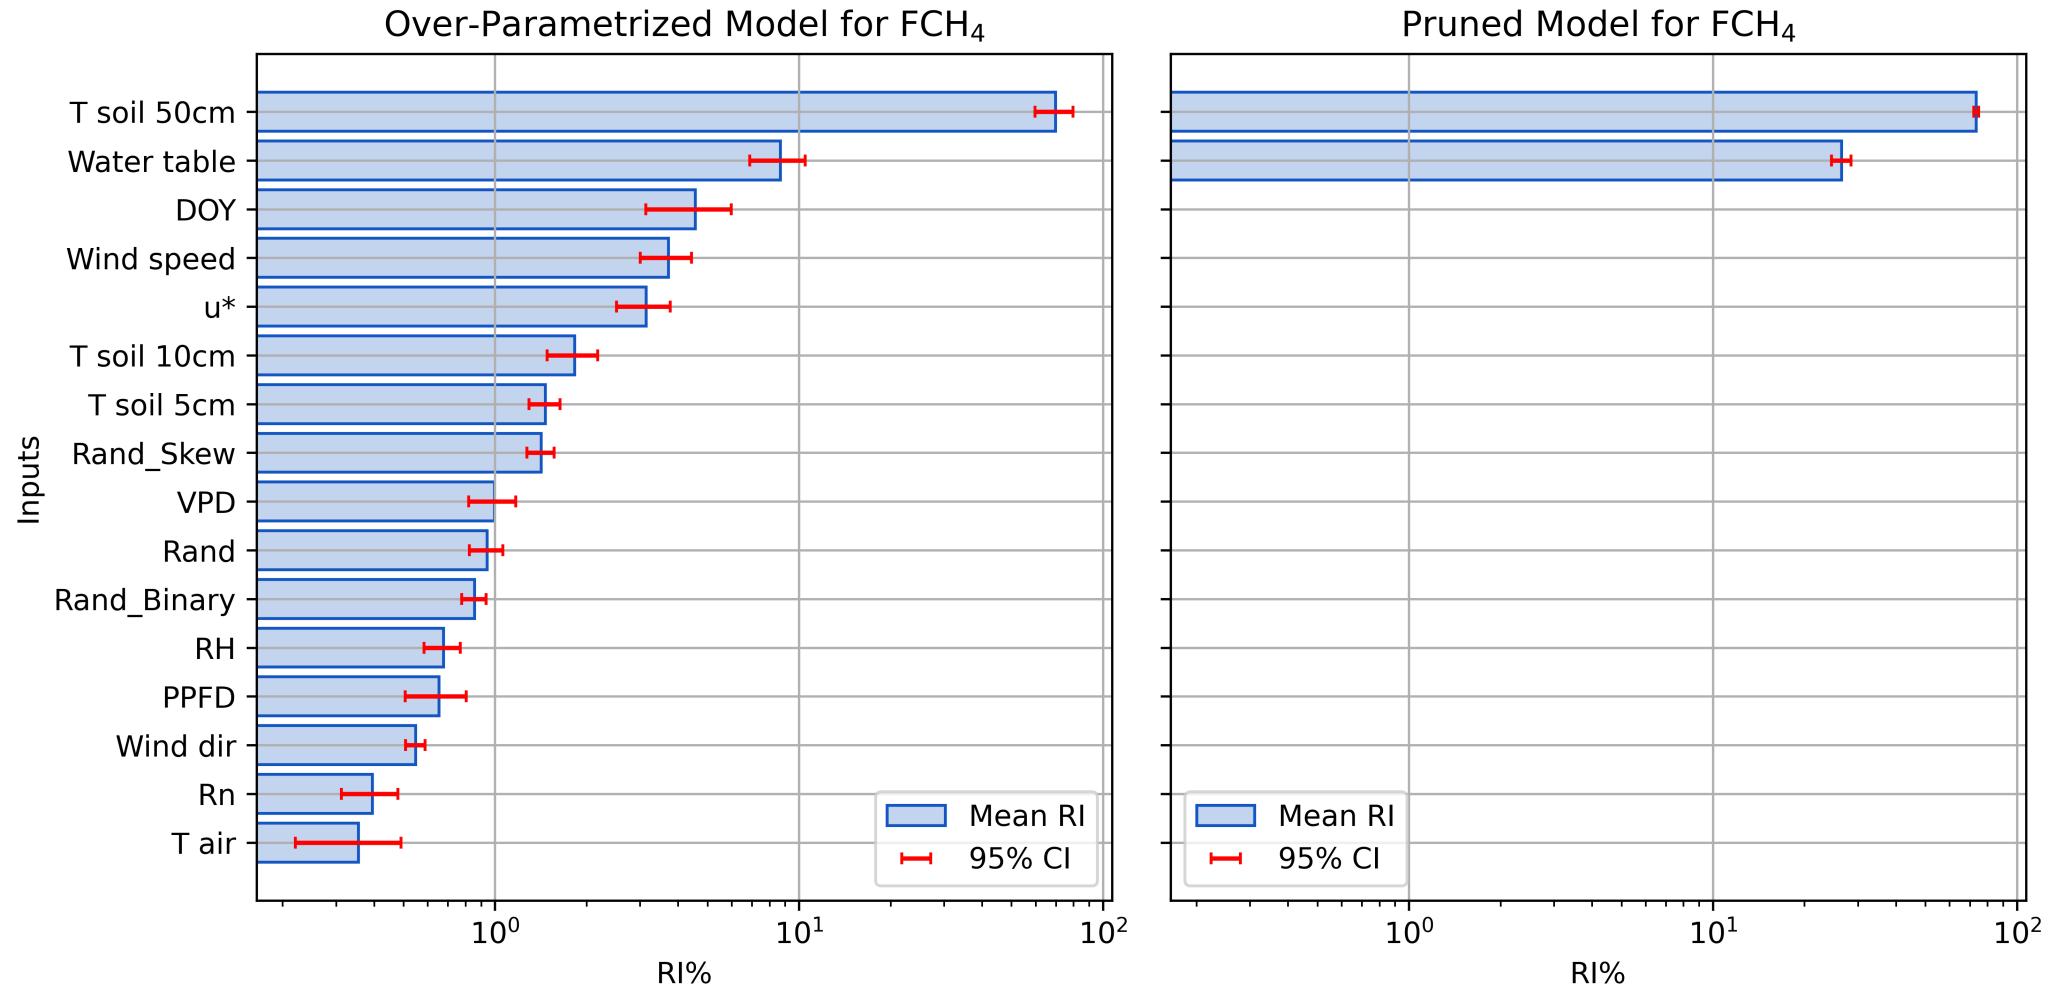
Model Performance FCO₂

Plot the model outputs and validation metrics calculated with the test data.

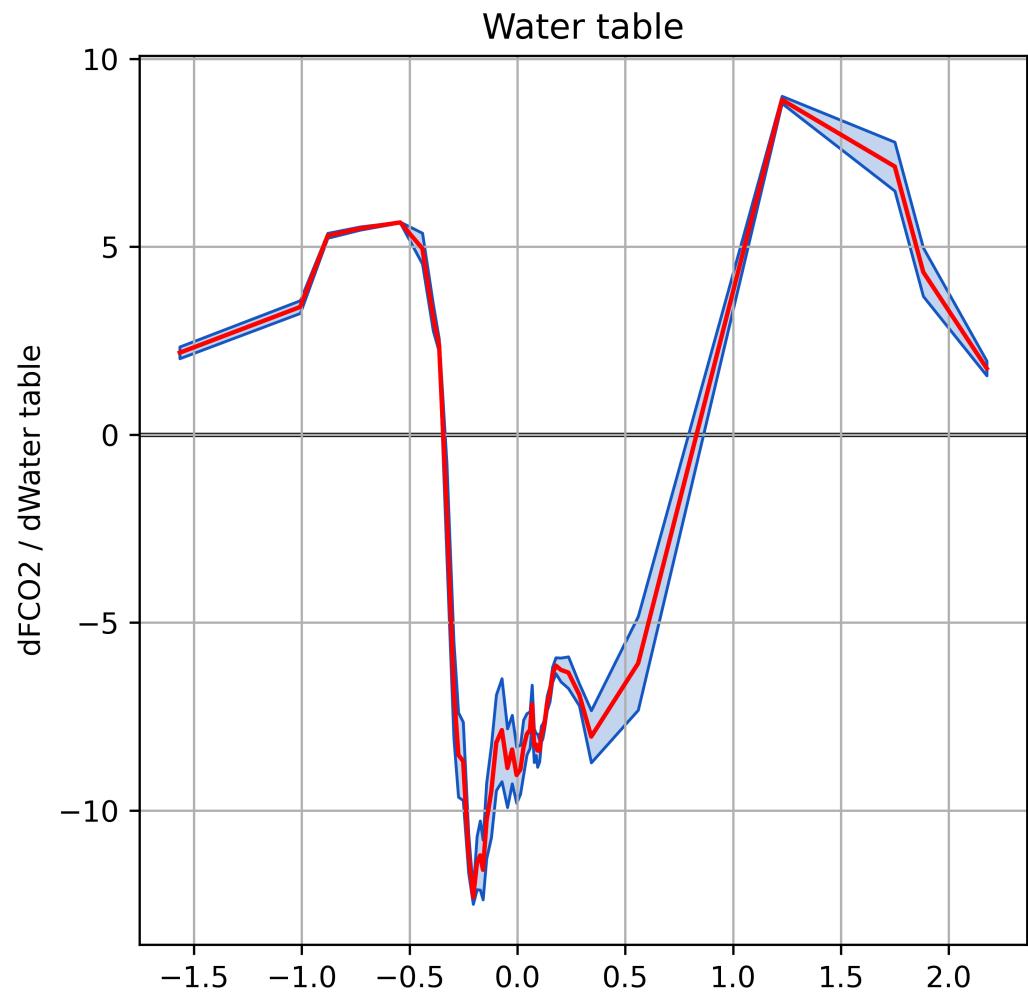
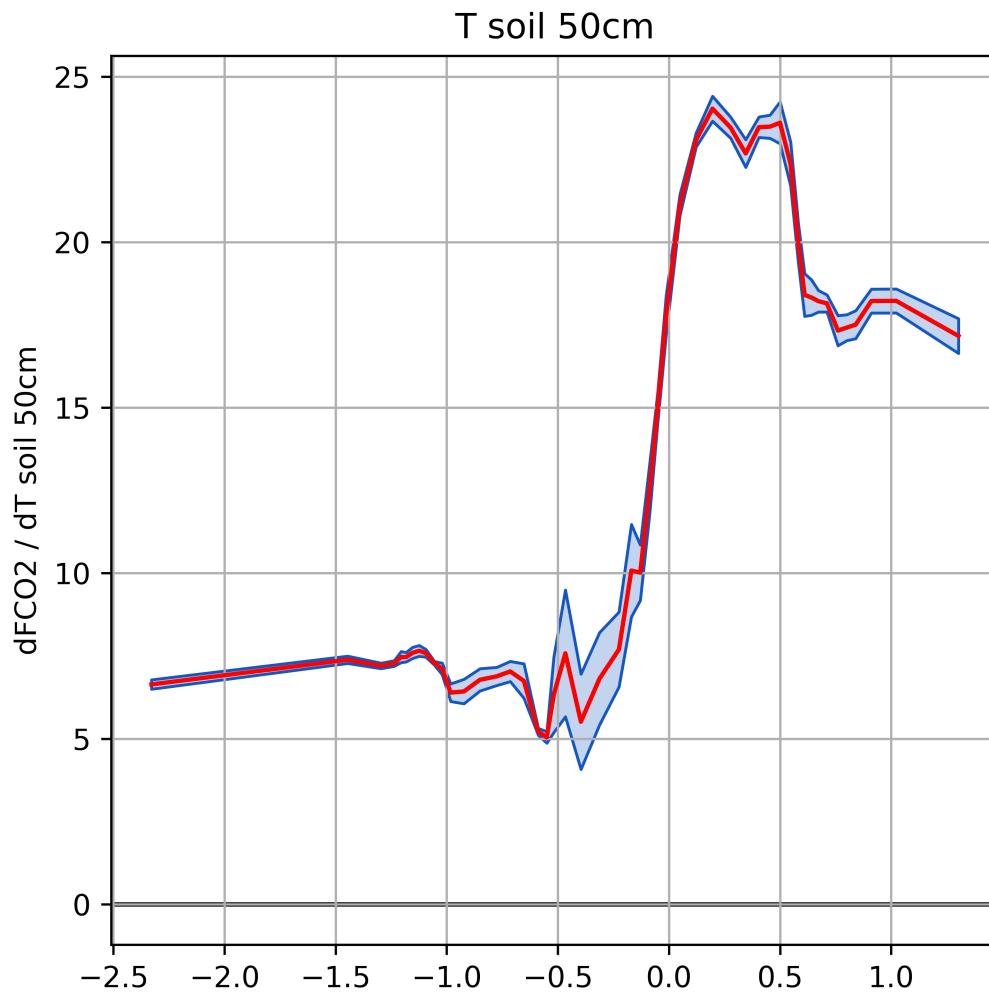
Metric	Score
RMSE	$0.45 \mu\text{mol m}^{-2} \text{s}^{-1}$
r ²	0.89



Before and After Pruning FCH₄



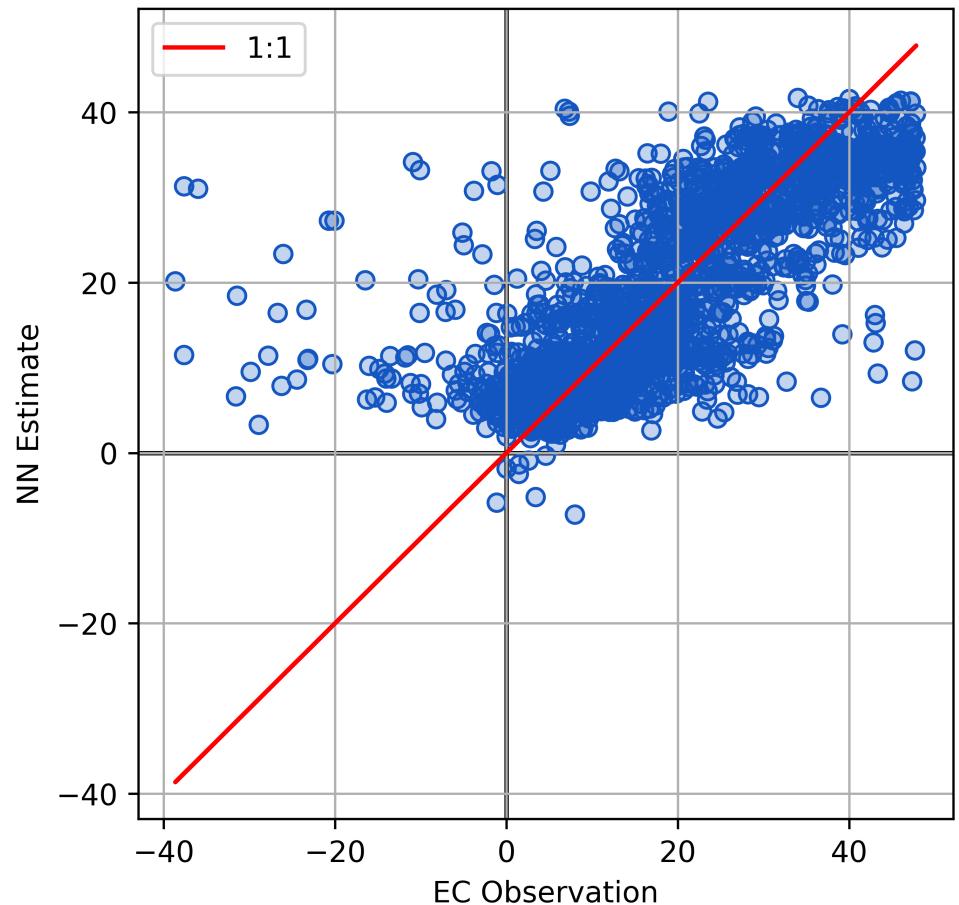
Normalized Derivatives of FCH₄



Model Performance FCH₄

Plot the model outputs and validation metrics calculated with the test data.

Metric	Score
RMSE	$8.58 \mu\text{mol m}^{-2} \text{s}^{-1}$
r ²	0.63



Next Steps

- Custom NN architecture: Separating input layers may allow us partition fluxes.
 - e.g., FCO_2 into GPP and ER
- Flux footprints: map response to spatial heterogeneity
- Upscaling: in space and time
- u^* filtering: partial derivatives could identify u^* thresholds
- Compare to process based models (e.g., CLASSIC)

Thank You

Questions?